

StereoQueerEval

Stereotype and Hate Speech Identification Targeting the LGBTQIA+ Community

WHY?

The exponential rise of social media has revolutionized online communication but has also intensified the circulation of harmful content, including hate speech, implicit discrimination, and deeply rooted stereotypes targeting LGBTQIA+ individuals. These stereotypes (often linguistically subtle) reinforce structural inequalities and harm queer communities globally.

Existing tasks such as LT-EDI (Kumaresan et al., 2025) and HODI (Nozza et al., 2023), along with the important work of the Queer in AI network, have advanced research on LGBTQIA+-related hate speech. However, both HODI and LT-EDI focus on homotransphobia detection: HODI is based on Italian tweets, while LT-EDI covers YouTube comments in English and several South Asian languages, with substantial imbalance across languages. In contrast, **our work aims to provide a balanced dataset that explicitly targets stereotypes about LGBTQIA+ people across multiple identities within the queer spectrum.**

WHAT'S NEW?

This edition builds on the expertise of three organizers previously involved in HODI (Nozza et al., 2023) and advances the state of the art in LGBTQIA+-related content detection across several dimensions.

- First, it explicitly addresses **stereotype identification**, making it possible to capture implicit biases and culturally embedded stereotypes about LGBTQIA+ people, including stereotypical content that may not be overtly hateful but still reinforces marginalization.
- Second, it captures the **implicit nature of hateful language** more broadly, extending the analysis beyond explicit hatred to include subtler forms of harmful expression.
- Third, it **considers multiple identities across the queer spectrum**, enabling a more fine-grained analysis than approaches that treat the LGBTQIA+ community as a single target category.
- Finally, it adopts a multilingual perspective with a multilingual dataset in **English, Italian, Dutch, and Farsi**.

DATA SOURCE

The task is to classify a YouTube user's comment, given the following contextual information:

- YouTube title
- Description of the video

This setting encourages the development of **context-aware systems**.

TEAM

Lead Organizer:

Alessandra T. Cignarella (Ghent University)

Co-Organizers:

Marta Marchiori Manerba (Turin University)
Greta Damo (Université Côte d'Azur, Inria)
Marzieh Abdolmaleki (Ghent University)

Advisory Organizers:

Debora Nozza (Bocconi University)
Els Lefever (Ghent University)

Language Support and Data Validation:

Joke Daems (Ghent University)
Joni Kruijsbergen (Ghent University)



SUBTASKS

A) Stereotype Identification: Binary stereotype detection (present/absent).

B) Hate Speech Type Classification: Multiclass hate speech detection, distinguishing between explicit hate, implicit hate, and non-hateful/supportive.

C) Target & LGBTQIA+ Identity Identification: Recognition of the target type (individual/group) and specific LGBTQIA+ identity (L, G, B, T, Q, I, A, Q+, NB).

SETTINGS & EVALUATION

Participants may choose freely which subtasks to take part in.

- Subtask A (Stereotypes) is fully self-contained.
- Subtask B (Hate Speech) is fully self-contained.
- Subtask C is conceptually designed as a cascade task following Subtask B, but we will not enforce any compulsory dependency.

Participants can therefore submit to any combination of subtasks (e.g., only A, only B, A+B, A+B+C, only C, B+C, etc...). They may also participate in any language or in any combination of languages.

We will produce one ranking per language for each subtask, resulting in 12 rankings in total (4 languages × 3 subtasks). Rankings will be based on **Macro F1**, following the methodology used in BRIGHTER 2025 (Muhammad et al., 2025) and POLAR 2026 (Naseem et al., 2026).

LANGUAGES

The task includes **four languages** with varying resources and sociolinguistic contexts:

- **English** → high-resource, broad online discourse, comparison with other shared tasks
- **Italian** → continuation of HODI 2023, mid-resource, European context, enabling intra-task comparison
- **Dutch** → mid-resource, European context, pioneering research queer-related resources
- **Farsi** → low-resource, culturally distinct, with underexplored queer-related discourse

We encourage participants to explore **multilingual transfer, cross-lingual learning**, low-resource approaches, and culturally grounded analyses.

INTEREST and PARTICIPANTS

Recent shared tasks at EVALITA showed a clear interest in **challenges focused on hate speech, discrimination, and LGBTQIA+-related harms**. HODI, the task on Homo- and Transphobia Detection at EVALITA 2023 collected around 30 submissions, as well as Multipride at EVALITA 2026 (on LGBTQIA+ slur reclamation) with 60 submissions, confirming engagement from the Italian NLP community.

Internationally, this trend is consistent: at SemEval the EDOS 2023 task, focused on sexism online, attracted 120 submissions for Task A, 70 for Task B and 50 for Task C. Similarly, the task POLAR 2026, about heated/polarized online debates targeting vulnerable groups (among which also LGBTQIA+ targets), received about 55 submissions for Task A and about 30 for Task B.

These participation levels show that socially relevant tasks draw substantial attention.